

Measuring the Defenders: A Layer-Aware, Framework-Mapped Benchmark for Model Context Protocol Security Proxies

Author: Gowthaman Arumugam, Independent Researcher (agowthaman90@gmail.com) **Version:** 0.3 · 2026-07-14 · *preprint, not peer-reviewed* **Artifacts:** mcp-defense-bench (benchmark) · mcp-bastion (reference proxy) — see Availability **Changes in v0.3:** two newly-published (June 2026) attacks added as vectors — mid-session tool injection (MSTI) and multi-tool split poisoning (ShareLock) — bringing the rubric to 24 vectors / 33 cases; the crosswalk gains a third U.S.-government framework column, the NSA MCP Security guidance (May 2026); mcp-bastion adds cross-tool correlation and now detects ShareLock (coverage 34% → 38%); and we add a tool-class taxonomy separating content-scanning proxies (scored) from access-control/isolation gateways (tracked).

Abstract

The Model Context Protocol (MCP) has become a common way to connect large-language-model (LLM) agents to external tools and data, and a substantial 2025–2026 literature now documents its attack surface. Existing security benchmarks for MCP measure how well an **LLM agent resists** attacks. None, to our knowledge, measure how well a **defensive proxy or gateway covers** that attack surface, and none map their results to the governance frameworks organizations actually use — the NIST AI Risk Management Framework (AI RMF) and the OWASP Top 10 for LLM and Agentic Applications. We present two artifacts to fill this gap: (1) a **threat-control crosswalk** of 24 MCP attack vectors mapped to architectural layer, STRIDE class, NIST AI RMF function, OWASP category, and the NSA MCP Security guidance; and (2) an open, vendor-neutral **defense-side benchmark** that runs a candidate MCP security proxy against a corpus of attack fixtures with matched benign controls, and scores its coverage using the crosswalk as a rubric. Each tool is driven through its *real* code, not a mock. We evaluate three open-source proxies (mcp-bastion, mcp-firewall, pipelock) plus a null baseline over 33 test cases. The strongest proxy reaches **38% weighted coverage (15 of 24 vectors)** — a proxy’s realistic ceiling — and even the union of all three tools leaves **9 of 24 vectors covered by none**: the entire registry/supply-chain surface, OS-isolation-dependent vectors, and several semantic vectors. All results carry zero false positives against matched benign controls. We argue this is direct, quantified evidence that a runtime proxy alone is insufficient and that defense-in-depth spanning *distinct mechanism classes* (proxy, cryptographic attestation, OS isolation) is required. We also report two methodological findings we observed first-hand: benchmark scores are sensitive to how attack fixtures are encoded (a fairness hazard), and the benchmark can *guide* a defender’s development — over the study, mcp-bastion’s coverage rose from 9% to 38% as gaps it surfaced were fixed and re-verified — including two attacks first published in June 2026, one of which (ShareLock) no measured tool covered until a cross-tool correlation check was added. An evasion-robustness dimension makes the encoding-sensitivity point concrete: obfuscated variants (zero-width, homoglyph, base64) of known attacks are caught by *different* tools, and no single tool survives all three. A matched benign-control design and a “measure only what the tool inspects at runtime” integrity rule keep the results honest.

1. Introduction

MCP standardizes how an LLM agent discovers and calls external “tools” exposed by “servers.” Its rapid adoption across consumer and enterprise agents has been accompanied, in 2025–2026, by a dense body of security research: systematization-of-knowledge papers, formal threat taxonomies, attack benchmarks, and cryptographic attestation proposals (Section 2). This work establishes *what can go wrong*.

A practical question remains under-addressed: **given a defensive tool that sits in front of MCP servers — a proxy, gateway, or scanner — how much of the known attack surface does it actually cover?** Answering it well requires three things that, individually, exist in the literature but have not been combined:

1. a **shared taxonomy** of MCP attack vectors;
2. a mapping from those vectors to the **governance frameworks** (NIST AI RMF, OWASP) that security and procurement teams reason in; and
3. an **executable, reproducible** way to test a defender and report coverage.

We contribute all three, plus an empirical study. Concretely:

- **A threat–control crosswalk** (Section 3): 24 MCP attack vectors, each mapped to one or more of six architectural layers, a STRIDE class, NIST AI RMF functions, NSA MCP Security recommendation areas, and OWASP LLM-2025 and Agentic-2026 categories. Published as open, versioned, machine-readable data.
- **A defense-side benchmark** (Section 4): for each vector, one or more attack **fixtures** paired with a near-identical **benign control**; a small **adapter** per tool that drives the tool’s real detection code; and a scorer that reports coverage per layer and per framework.
- **An empirical evaluation** (Sections 5–6) of three open-source proxies and a null baseline, yielding the complementarity finding above.
- **A candid methodology discussion** (Sections 7–8): the corpus-encoding sensitivity we observed, the matched-control and runtime-integrity safeguards, and the limits of the study.

Our aim is not to rank products but to map coverage, expose gaps, and provide reusable infrastructure.

2. Related Work

All 2603/2604-series works below are recent, non-peer-reviewed preprints whose taxonomies are their authors’ own constructs; we cite them as *proposed by*, not as established consensus. Coverage statistics quoted from any single paper reflect that paper’s surveyed set.

Threat taxonomies. Recent work systematizes MCP threats through several lenses: an SoK separating adversarial security threats (indirect prompt injection, tool poisoning) from epistemic safety hazards across the Resources/Prompts/Tools primitives [SoK-2512.08290]; a formal framework of 7 threat categories and 23 attack vectors grounded in a large tool corpus [Formal-2604.05969]; a STRIDE/DREAD model over five MCP components that identifies tool poisoning as the most impactful client-side vulnerability [STRIDE-2603.22489]; and a defense-placement taxonomy classifying defenses by the layer that enforces them, which finds protection “uneven and predominantly tool-centric” with weak transport, host-orchestration, and supply-chain coverage [MCP-DPT-2604.07551]. Our crosswalk reuses these vector families but adds the framework mapping none of them provides.

Attack-side benchmarks. MSB [MSB-2510.15994] (12 attacks, ~2,000 instances, real MCP execution) measures the resistance of the **LLM agent** across the tool-use pipeline; MCPTox [MCPTox-2508.14925] scopes to tool poisoning; AgentDefense-Bench aggregates datasets into MCP-JSON-RPC test cases; MCP-Bench [MCPBench-2508.20453] measures agent *capability*, not security. All measure the agent (or the model), not a **defensive proxy**. Ours measures the defender, on the full mapped surface, and reports coverage crosswalked to NIST/OWASP — which we could not find in any surveyed benchmark.

Attestation and provenance. ETDI [ETDI-2506.01333] proposes signed, versioned tool definitions; a “Trustworthy MCP Registry” blueprint [TrustReg-FI2026] composes existing standards (well-known URIs, Sigstore, JSON Canonicalization + JWS) for provenance and runtime integrity, and notes the official MCP registry remains an unverified pointer architecture. These are complementary to our benchmark; a future attestation-conformance suite (Section 9) would test them.

Scanners and gateways. Practitioner tools include Invariant Labs’ MCP-Scan (since acquired by Snyk and now a cloud-gated scanner), Lasso’s and EnkryptAI’s MCP gateways, Docker’s isolation-based MCP gateway, and the proxies we evaluate here. Our benchmark scores locally-deterministic defenders and treats cloud-model defenders as observable-but-not-reproducibly-scoreable (Section 4.4).

3. The Threat–Control Crosswalk

The crosswalk is the benchmark’s rubric and a standalone artifact. It enumerates **24 MCP attack vectors** consolidated from the taxonomies in Section 2, and maps each to:

- **Architectural layer** (one of: host-orchestration, client, transport, server, tool, registry-supply-chain) — following the defense-placement view of [MCP-DPT-2604.07551];
- **STRIDE** class;
- **NIST AI RMF** function(s): Govern, Map, Measure, Manage;
- **NSA MCP Security** recommendation area (Authentication & Access Control, Monitoring & Logging, Data Classification & Segmentation, Runtime Controls) — from the NSA AI Security Center guidance [NSA-MCP-2026];
- **OWASP Top 10 for LLM Applications (2025)** category (LLM01–LLM10); and
- **OWASP Top 10 for Agentic Applications (2026)** category (ASI01–ASI10).

Table 1 lists the vectors and their primary layer. The full mapping, including STRIDE and all framework codes, is published as machine-readable data (`rubric/crosswalk.json`, CC-BY-4.0).

Table 1. The 24 vectors (primary layer).

#	Vector	Layer	#	Vector	Layer
1	Tool poisoning	tool	13	Sandbox escape	server
2	Tool shadowing / name collision	client	14	Schema / validation bypass	server
3	Rug pull (definition mutation)	tool	15	Man-in-the-middle (transport)	transport
4	Out-of-scope parameter injection	tool	16	DNS rebinding (local servers)	transport
5	Prompt injection via tool results	tool	17	Server impersonation	registry-supply-chain
6	Indirect / retrieval injection	tool	18	Excessive permission / escalation	host-orchestration
7	Cross-tool exfiltration (confused deputy)	client	19	Credential / token theft	host-orchestration
8	Tool-transfer / cross-server chaining	host-orchestration	20	Consent fatigue / over-broad grants	client
9	False-error escalation	tool	21	Command injection	server
10	Package / name squatting	registry-supply-chain	22	System-prompt / context leakage	client
11	Supply-chain poisoning (provenance gap)	registry-supply-chain	23	Mid-session tool injection (MSTI)	client

#	Vector	Layer	#	Vector	Layer
12	Configuration drift	server	24	Multi-tool split poisoning (ShareLock)	tool

The mappings are indicative and reviewable, not a certification — we invite community correction, which is itself a goal of publishing the rubric openly. The OWASP Agentic categories used are the official 2026 set: ASI01 Agent Goal Hijack, ASI02 Tool Misuse and Exploitation, ASI03 Identity and Privilege Abuse, ASI04 Agentic Supply Chain Vulnerabilities, ASI05 Unexpected Code Execution (RCE), ASI06 Memory & Context Poisoning, ASI07 Insecure Inter-Agent Communication, ASI08 Cascading Failures, ASI09 Human-Agent Trust Exploitation, ASI10 Rogue Agents.

4. Benchmark Methodology

4.1 Design

For each vector the corpus contains one or more **test cases**. A test case is a static JSON object with: a **malicious fixture** (the attack, as data — a tool definition, tool call, tool result, registry entry, or scenario), a near-identical **benign control**, the **expected** defender behavior (**detect/enforce**), and a **provenance** note. Fixtures are static data, never live exploits; attacker hosts use reserved names (**attacker.example**; the cloud-metadata address `169.254.169.254` appears only as an SSRF *target* in data, never contacted).

4.2 Adapters and the runtime-integrity rule

Each evaluated tool provides a thin **adapter** exposing `assess(input) → {detect, enforce}`. The runner calls it once on the malicious fixture and once on the benign control. Adapters drive the tool’s **real** code (e.g., importing the proxy’s detection functions, or invoking its CLI/SDK with a committed configuration).

A central rule keeps scores honest:

Report only what the tool actually inspects at runtime — never what its rules *could* match if pointed at data the tool never sees.

For example, a proxy that inspects tool *definitions* but not tool *results* returns “not detected” for result-borne attacks, even if its text rules would match the result string in isolation. This measures deployed behavior, not theoretical capability.

4.3 Scoring

Per test case: if the malicious fixture is not flagged, the vector scores **none**; if it is flagged but the benign control is *also* flagged, the vector scores **none** and is marked a **false positive**; if the malicious fixture is flagged and the benign control is clean, the vector scores **detect** (warn) or **enforce** (block). A vector’s level is the best across its test cases. Weighted coverage assigns `enforce = 1.0`, `detect = 0.5`, `none = 0`, over 24 vectors. The matched benign control is essential: without it, a tool that flags everything would score perfectly while being useless.

4.4 Reproducibility rule

A defender is **scored** only if its detection logic runs **locally and deterministically**. Tools that delegate detection to a proprietary cloud model can be *observed* but not reproducibly *scored*, since their verdicts depend on a black box that can change between runs. This excludes, e.g., the cloud-gated Snyk agent-scan; it includes the three proxies below.

5. Experimental Setup

We evaluate four adapters over a 33-case corpus (24 base cases + 6 realistic “v2” cases for fairness + 3 “v3” evasion-robustness cases; see Section 7):

Tool	Class	License	How driven
null-baseline	control	—	returns “not detected” for all inputs
mcp-bastion	runtime proxy	Apache-2.0	imports real detection: <code>scanTool</code> , <code>scanText</code> , <code>scanToolSet</code> (cross-tool correlation), <code>validateArguments</code> , <code>checkRequestedScopes</code> , <code>checkTransportSecurity</code> , <code>checkRequestOrigin</code> , <code>hashToolDefinition</code>
mcp-firewall	runtime proxy (Python)	AGPL-3.0	real SDK (<code>Gateway.check / scan_response</code>) with a committed config
pipelock	egress firewall (Go)	Apache-2.0 core	real scanners via <code>explain --json</code> (URL/egress) and <code>mcp scan</code> (MCP-response injection), both offline and deterministic

For AGPL/other licensed tools we *run* the tool to benchmark it and report scores; we do not redistribute their code. Each tool is pinned to a released version and configured with its own recommended/starter policy, committed to the repository for reproducibility.

6. Results

Table 2. Verified coverage (33 cases; weighted over 24 vectors).

Tool	Class	Weighted coverage	enforce	detect	none	False positives
mcp-bastion	runtime proxy	38% (9.0/24)	3	12	9	0 / 33
mcp-firewall	runtime proxy	13% (3.0/24)	3	0	21	0 / 33
pipelock	egress firewall	10% (2.5/24)	1	3	20	0 / 33
null-baseline	control	0%	0	0	24	0 / 33

A proxy’s ceiling is partial. mcp-bastion, the broadest tool, spans the definition, response, argument, transport, and egress/least-privilege layers (poisoning, shadowing, rug-pull; response and retrieval injection and result-borne leakage; schema/parameter validation; plaintext-transport and DNS-rebinding defense; sensitive-argument and over-broad-scope detection). mcp-firewall and pipelock add *enforcement* (deny/redact/SSRF-block) on the call and egress layers where bastion only warns; pipelock additionally detects response-borne prompt injection via its MCP-response scanner.

Two newly-published attacks (June 2026). We added mid-session tool injection (MSTI, [WebMCP-MSTI-2606.06387]) and multi-tool split poisoning (ShareLock, [ShareLock-2606.27027]) as vectors 23–24. MSTI — a tool that mutates its behavior after first approval, mid-session — is caught by mcp-bastion’s definition pinning, which re-hashes the tool on each use and *enforces* on a post-pin change; neither other tool covers it. ShareLock — a payload split across several tools so no single description looks malicious — was covered by **no** measured tool until we added **cross-tool correlation** to mcp-bastion (Section 7): because a proxy sees a server’s whole tool set, it scans the combined descriptions and flags coordinated **share/checksum**-style staging metadata across tools. It now *detects* ShareLock; the other two tools, which scan one tool at a time, still miss it. This is a staging-signal heuristic, not a cryptographic defeat of threshold secret-sharing (Section 8).

Defense-in-depth across mechanism classes is required. Even with the strongest proxy at 38%, and across all three tools, **15 of 24 vectors are covered by at least one tool; 9 are covered by none** — the registry/supply-chain surface (package squatting, provenance-gap supply-chain poisoning, server impersonation), OS-isolation-dependent vectors (sandbox escape), and semantic vectors (tool-transfer, false-error escalation, configuration drift, command injection, consent fatigue). These are not addressable by *any* runtime proxy; they require different mechanisms — cryptographic attestation, OS sandboxing, and client-side controls. A proxy is necessary but not sufficient.

Evasion robustness. The corpus includes three obfuscated variants of known attacks — zero-width / bidirectional-control characters, Cyrillic homoglyph substitution, and base64-wrapped payloads. The tools resist *different* evasions: mcp-bastion catches the zero-width/bidi variant (via a hidden-character rule) but misses the others; pipelock catches the homoglyph and base64 variants (via a normalization pipeline) but misses zero-width; mcp-firewall catches none. No single tool survives all three, so the defense-in-depth conclusion holds at the evasion layer as well as the vector layer.

Zero false positives. No tool flagged any benign control, across all 33 cases and every feature iteration.

7. Discussion

Benchmark-guided improvement. The benchmark initially scored mcp-bastion at 9%, covering only the tool-definition layer. It then guided five rounds of proxy-native hardening, each re-measured at zero false positives: response scanning (9%→18%: response/retrieval injection, result-borne leakage), argument/schema validation (18%→23%: parameter smuggling, validation bypass), transport hardening (23%→30%: plaintext-transport warning and a DNS-rebinding Origin check that *enforces* by rejecting cross-origin requests to a loopback listener), sensitive-argument plus least-privilege scanning (30%→34%: cross-tool exfiltration, overbroad scopes), and cross-tool correlation (34%→38%: ShareLock split poisoning, plus MSTI caught by existing pinning). The benchmark did not just measure the tool; it identified concrete, layer-specific gaps and verified each fix — and then defined the proxy’s ceiling, since the remaining 9 vectors are architecturally outside a proxy’s reach.

A living benchmark tracks new attacks. Vectors 23–24 were published (June 2026) *after* the v0.2 rubric was frozen; adding them exercised the benchmark’s intended lifecycle. Each was verified against its primary source, encoded as a matched malicious/benign pair, and re-run against every tool — which immediately exposed that ShareLock was covered by none, driving the cross-tool-correlation fix above. A benchmark that only measured a fixed attack set would have missed the gap entirely; the value is in re-measuring as the threat surface moves.

Tool-class taxonomy (what we score vs. what we track). Not every MCP defense is a content-scanning proxy, and scoring all of them on this rubric would be misleading. We separate two classes. *Content-scanning proxies* (mcp-bastion, mcp-firewall, pipelock) inspect tool definitions, calls, or results at runtime; the rubric measures exactly what they inspect, so we **score** them. *Access-control and isolation gateways* (e.g. MCPX/Lunar, IBM ContextForge, Docker’s isolation gateway) enforce identity, policy, and sandboxing rather than scanning content; the rubric’s content-borne vectors would score them near zero not because they are weak but because they operate on a different axis. We **track** these in the landscape (Section 2) without assigning a score, and note that a class-appropriate enforcement-and-isolation rubric is future work. Reporting a misleading number would be worse than reporting none.

Corpus-encoding sensitivity (a fairness hazard). We observed that scores are sensitive to how an attack is *encoded*. An early corpus expressed exfiltration with sanitized placeholder text; against it, mcp-firewall scored 5%. Adding realistic encodings (an AWS-key-shaped secret in a tool result; an exfiltration POST to the cloud-metadata address) raised it to 13% — not by changing the tool, but by probing capabilities the first corpus never exercised. We therefore treat *tool-neutral, realistic, multi-encoding* fixtures as a fairness requirement, and we caution that **absolute totals are corpus- dependent; the per-vector matrix (which vectors a tool covers) is the more reliable read.**

Adapter faithfulness. The matched benign control also guards against adapter over-reach. In wiring pipelock (an egress firewall), an initial adapter synthesized an egress URL from an *inbound* local host, causing pipelock to “block” both the malicious and benign case — surfaced immediately as a false positive and corrected by restricting URL synthesis to outbound endpoints.

8. Limitations and Threats to Validity

- **Taxonomy is a synthesis, not a standard.** The 24-vector set consolidates several proposed taxonomies; it is one reviewable synthesis. Framework mappings are indicative, single-author, and pending a second-reviewer pass; OWASP Agentic ASI05–ASI10 labels are unconfirmed.
- **Small, seeded corpus.** Most vectors have 1–2 fixtures. Absolute coverage numbers should be read as lower bounds and are corpus-dependent (Section 7).
- **The ShareLock check is a heuristic, not a proof.** Cross-tool correlation flags the *staging signal* — coordinated *share/checksum*-style metadata across a tool set — that the published attack uses. It is not a cryptographic defeat of threshold secret-sharing; a sufficiently disguised split payload that avoids that signal can still evade it. We report it as **detect**, not **enforce**.
- **Partial tool interfaces.** For pipelock we drive its offline URL/egress scanner and its MCP-response injection scanner, but not its poisoned-tool-description check or its full DLP/tool-call HTTP eval endpoint; its number therefore remains a lower bound on its total capability.
- **Configuration dependence.** Results depend on each tool’s configured policy; we commit the exact configs, but different policies would yield different numbers.
- **Scope and tool class.** We score locally-deterministic content-scanning proxies; cloud-model defenders and access-control/isolation gateways are out of scope for scoring (Sections 4.4, 7) — the latter would score misleadingly low on a content-borne rubric and are tracked, not scored.

9. Future Work

Expand the corpus to multiple tool-neutral fixtures per vector and deepen the evasion-robustness set; drive pipelock’s remaining interfaces (poisoned-description and DLP/tool-call eval endpoint) and add further defenders (including a fully-local mode of cloud-gated gateways, if exposed); publish a public leaderboard; complete the second-reviewer pass on framework mappings and confirm ASI labels; and build an **attestation-conformance** suite to test tool-integrity/provenance schemes (ETDI, the Trustworthy Registry composition) as a coupled bench-plus-attestation artifact.

10. Conclusion

We introduced a layer-aware, framework-mapped crosswalk of 24 MCP attack vectors — mapped to NIST AI RMF, the OWASP LLM and Agentic Top 10s, STRIDE, and the NSA MCP Security guidance — and an open, vendor-neutral benchmark that measures how much of that surface a defensive proxy actually covers, driving each tool through its real code. The strongest of three open-source proxies reaches 38% coverage, and 9 of 24 vectors are undefended by any measured tool — the registry/supply-chain, OS-isolation, and semantic vectors that lie outside a proxy’s reach. The evidence argues for layered defense across distinct mechanism classes, and the benchmark demonstrably guided one proxy from 9% to 38% at zero false positives — including a cross-tool-correlation check added in response to a June-2026 attack that no measured tool had covered. It provides a reusable, honest instrument — with matched benign controls, a runtime-integrity rule, and a reproducibility rule — for the community to measure MCP defenders and track their progress.

Availability

- **Benchmark (mcp-defense-bench):** rubric, corpus, adapters, runner, leaderboard. Rubric/data: CC-BY-4.0; harness: Apache-2.0.
- **Reference proxy (mcp-bastion):** Apache-2.0; published on npm.
- Evaluated third-party tools are used under their own licenses (mcp-firewall: AGPL-3.0; pipelock: Apache-2.0 core) and are not redistributed here.
- **Preprint of this paper:** DOI 10.6084/m9.figshare.32978657
- **mcp-defense-bench:** <https://github.com/Gowthaman90/mcp-defense-bench> · archived on Zenodo, DOI 10.5281/zenodo.21346206
- **mcp-bastion:** <https://github.com/Gowthaman90/mcp-bastion> · npm: **mcp-bastion**

References

Titles, authors, and identifiers below were verified against primary sources on 2026-07-14. Each arXiv item resolves at <https://arxiv.org/abs/<id>>. The 2603/2604/2606-series works are recent, non-peer-reviewed preprints and may have updated versions; verify current version and page/DOI details at camera-ready time.

- [SoK-2512.08290] S. Gaire, S. Gyawali, S. Mishra, S. Niroula, D. Thakur, U. Yadav. “Systematization of Knowledge: Security and Safety in the Model Context Protocol Ecosystem.” arXiv:2512.08290, Dec. 2025.
- [Formal-2604.05969] N. Acharya, G. K. Gupta. “A Formal Security Framework for MCP-Based AI Agents: Threat Taxonomy, Verification Models, and Defense Mechanisms.” arXiv:2604.05969, Apr. 2026.
- [STRIDE-2603.22489] C. Huang, X. Huang, N. P. Tran, A. Milani Fard. “Model Context Protocol Threat Modeling and Analyzing Vulnerabilities to Prompt Injection with Tool Poisoning.” arXiv:2603.22489, Mar. 2026.
- [MCP-DPT-2604.07551] M. Rostamzadeh, S. Narula, N. Birhan, M. Ghasemigol, D. Takabi. “MCP-DPT: A Defense-Placement Taxonomy and Coverage Analysis for Model Context Protocol Security.” arXiv:2604.07551, Apr. 2026.
- [MSB-2510.15994] D. Zhang, Z. Li, X. Luo, X. Liu, P. Li, W. Xu. “MCP Security Bench (MSB): Benchmarking Attacks Against Model Context Protocol in LLM Agents.” arXiv:2510.15994, Oct. 2025 (ICLR 2026).
- [MCPTox-2508.14925] Z. Wang, Y. Gao, Y. Wang, S. Liu, H. Sun, H. Cheng, G. Shi, H. Du, X. Li. “MCPTox: A Benchmark for Tool Poisoning Attack on Real-World MCP Servers.” arXiv:2508.14925, Aug. 2025.
- [MCPBench-2508.20453] Z. Wang, Q. Chang, H. Patel, S. Biju, C.-E. Wu, Q. Liu, A. Ding, A. Reza-zadeh, A. Shah, Y. Bao, E. Siow. “MCP-Bench: Benchmarking Tool-Using LLM Agents with Complex Real-World Tasks via MCP Servers.” arXiv:2508.20453, Aug. 2025.
- [ETDI-2506.01333] M. Bhatt, V. S. Narajala, I. Habler. “ETDI: Mitigating Tool Squatting and Rug Pull Attacks in Model Context Protocol (MCP) by using OAuth-Enhanced Tool Definitions and Policy-Based Access Control.” arXiv:2506.01333, Jun. 2025.
- [TrustReg-FI2026] L. Mas, J. Vilaplana, J. Rius, R. Spaimoc, J. Mateo. “The Trustworthy Model Context Protocol (MCP) Registry: An Architectural Blueprint for Cryptographic Provenance and Runtime Integrity.” Future Internet, vol. 18, no. 5, art. 243, 2026. doi:10.3390/fi18050243.
- [WebMCP-MSTI-2606.06387] “WebMCP: Tool-Surface Poisoning and Mid-Session Tool Injection in Browser-Hosted Model Context Protocol.” arXiv:2606.06387, Jun. 2026. (Vector 23.)
- [ShareLock-2606.27027] “ShareLock: A Stealthy Multi-Tool Threshold Poisoning Attack Against the Model Context Protocol.” arXiv:2606.27027, Jun. 2026. (Vector 24.)
- [NIST-AI-RMF] National Institute of Standards and Technology (U.S. Dept. of Commerce). “AI Risk Management Framework (AI RMF 1.0).” NIST AI 100-1, 2023 (functions: Govern, Map, Measure, Manage). <https://www.nist.gov/itl/ai-risk-management-framework> · <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.100-1.pdf>

- [OWASP-LLM-2025] OWASP GenAI Security Project. “OWASP Top 10 for LLM Applications,” 2025 edition. <https://genai.owasp.org/llm-top-10/>
- [OWASP-ASI-2026] OWASP GenAI Security Project. “OWASP Top 10 for Agentic Applications,” 2026 (published Dec. 9, 2025; categories ASI01–ASI10). <https://genai.owasp.org/resource/owasp-top-10-for-agentic-applications-for-2026/>
- [NSA-MCP-2026] National Security Agency, AI Security Center. “MCP Security: Security Design Considerations for AI-Driven Automation” (Cybersecurity Information Sheet, U/OO/6030316-26), May 2026. Recommendation areas: Authentication & Access Control, Monitoring & Logging, Data Classification & Segmentation, Runtime Controls. https://www.nsa.gov/Portals/75/documents/Cybersecurity/CSI_MCP_SE
- [STRIDE] Microsoft. “STRIDE — Threat Modeling.” <https://learn.microsoft.com/en-us/azure/security/develop/threat-modeling-tool-threats>
- [AgentDefense-Bench] A. Sanna. “AgentDefense-Bench” (open-source dataset), GitHub.
- [Snyk-agent-scan] Snyk. “agent-scan” (formerly Invariant Labs MCP-Scan), GitHub / vendor tooling.